



This presentation is released under the terms of the
Creative Commons Attribution-Share Alike license.

You are free to reuse it and modify it as much as you want as long as:

- (1) you mention Séverin Lemaignan as being the original author,
- (2) you re-share your presentation under the same terms.

You can download the sources of this presentation here:

github.com/severin-lemaignan/lecture-hri-ml-for-hri

brl

UWE
Bristol

University
of the
West of
England



University of
BRISTOL

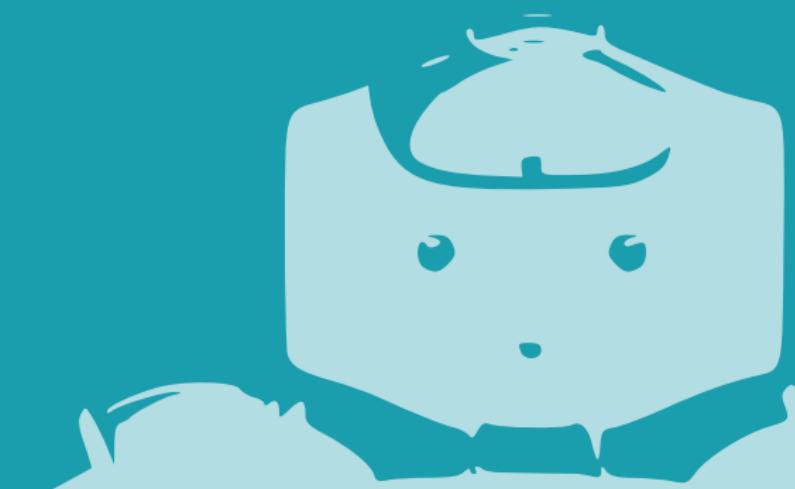
Human-Robot Interaction

Machine Learning for HRI

Séverin Lemaignan

Bristol Robotics Lab

University of the West of England



Emotions?
oo

Classification
oooooooooooooooooooo

Emotion classifier
oooooooooooo

ML for social behaviours generation
oooooooooooo

IN THIS LECTURE

- o Build your own emotion classifier from scratch!
- o Basic classifiers: kNNs, SVM, MLP
- o ML to teach a robot social behaviours

EMOTIONS?

PAUL EKMAN AND BASIC EMOTIONS

Paul Ekman, a psychologist, found that when shown facial expressions people across the world all recognised six **basic emotions**.



PAUL EKMAN AND BASIC EMOTIONS

Paul Ekman, a psychologist, found that when shown facial expressions people across the world all recognised six **basic emotions**.

Anger, disgust, fear, happiness, sadness and surprise.

PAUL EKMAN AND BASIC EMOTIONS

Paul Ekman, a psychologist, found that when shown facial expressions people across the world all recognised six **basic emotions**.

Anger, disgust, fear, happiness, sadness and surprise.

In 1990 Ekman extended his list of basic emotion to include other emotions: *Amusement, contempt, contentment, embarrassment, excitement, guilt, pride in achievement, relief, satisfaction, sensory pleasure, and shame*

Attention

Ekman's model does not adequately reflect the real-world complexity of emotions. Today, we favour continuous models like valence-arousal, or Pleasure-Arousal-Dominance (PAD), or discuss more general *facial expressions*.

CLASSIFICATION

Emotions?
oo

Classification
o●oooooooooooooooooooo

Emotion classifier
oooooooooooo

ML for social behaviours generation
oooooooooooo

CLASSIFICATION

Social signal processing often relies on **classification**.

- **Classification** is deciding on which **category** a new observation belongs to based on training data.
- The training data contains data and known categories.
- Classification is a **supervised** learning algorithm.

CLASSIFICATION

Social signal processing often relies on **classification**.

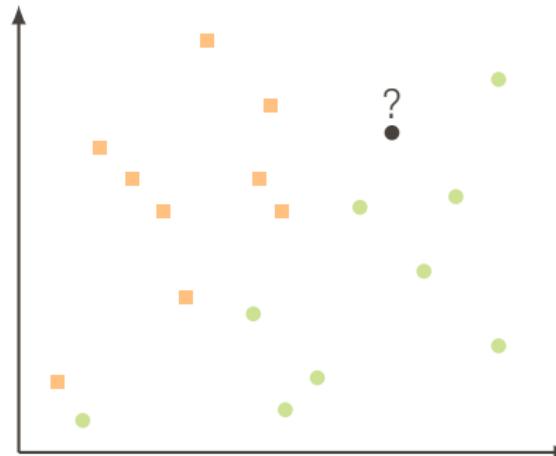
- **Classification** is deciding on which **category** a new observation belongs to based on training data.
- The training data contains data and known categories.
- Classification is a **supervised** learning algorithm.

Examples:

- An incoming tweet needs to be classified as being positive or negative.
- A radar ping of a flying objects needs to be classified as belonging to one of n possible planes.
- The prosody of speech needs to be classified as belonging to one of six basic categories of emotion (happy, sad, angry, bored, surprised, neutral).
- A gesture filmed through a camera needs to be classified as meaning stop, go, left or right.

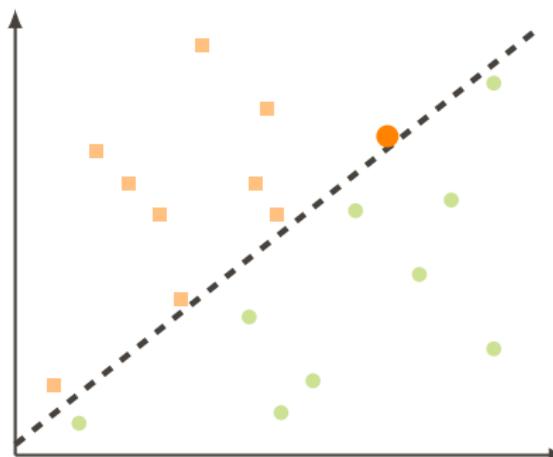
CLASSIFICATION: EXAMPLE

- o Two dimensional problem
- o Two categories, with training data for both categories
- o To which category does a new observation belong?



CLASSIFICATION: EXAMPLE

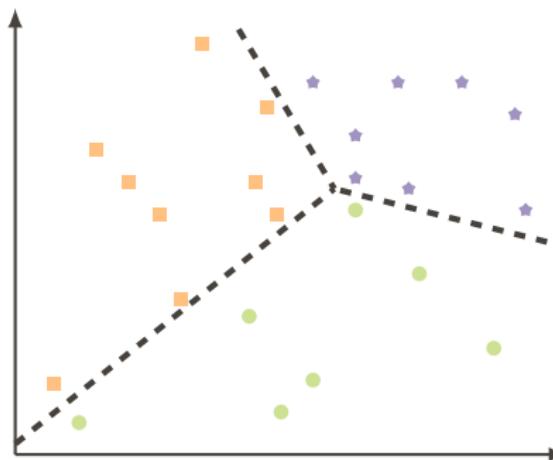
- Two dimensional problem
- Two categories, with training data for both categories
- To which category does a new observation belong?



When categories can be linearly separated, we have a very easy classification problem.

CLASSIFICATION: EXAMPLE

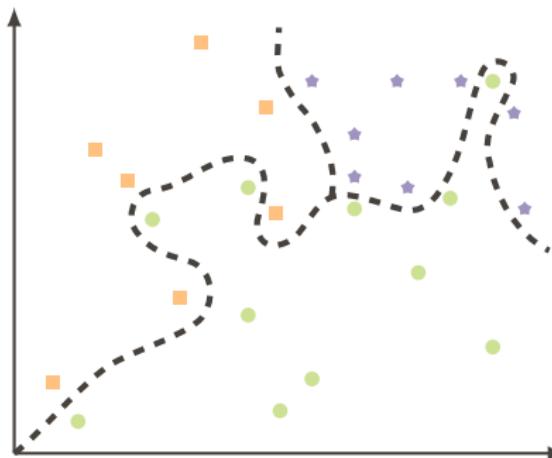
- Two dimensional problem
- Two categories, with training data for both categories
- To which category does a new observation belong?



Three or more can still be linearly separated.

CLASSIFICATION: EXAMPLE

- Two dimensional problem
- Two categories, with training data for both categories
- To which category does a new observation belong?



But what if categories are not linearly separable?

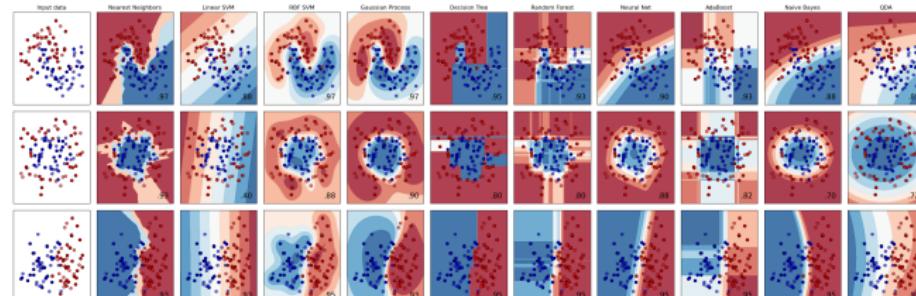
THREE CLASSIFICATION METHODS

Three methods will be explained here:

- k -nearest Neighbours (k NN)
- Support Vector Machines (SVM)
- Multi-layer perceptrons (MLP)

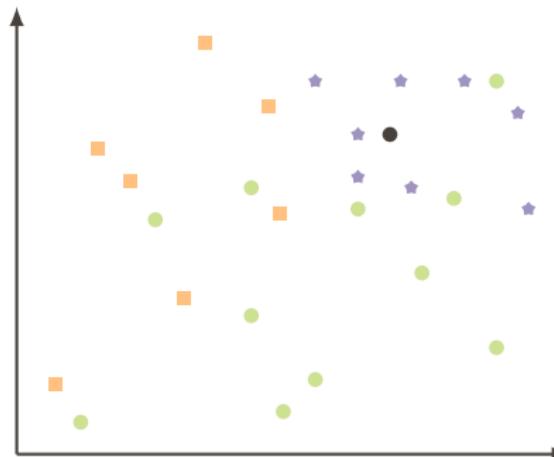
Many other classifiers are possible (Comparison of classifiers with scikit-learn):

- Decision trees/Random forest
- Bayes classifiers
- The large family of neural networks (MLP is the simplest one)
- ...



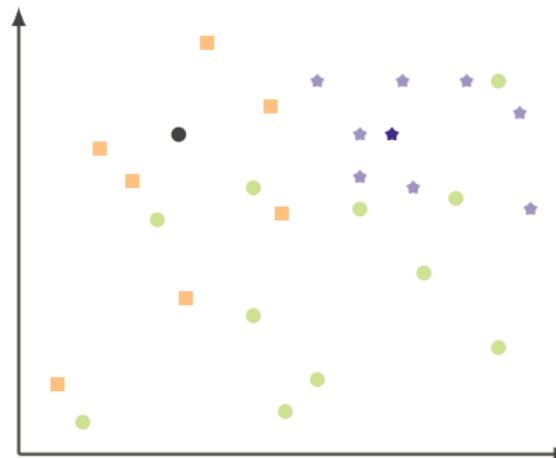
K-NEAREST NEIGHBOURS

When an observation comes in, calculate the distance to the k nearest neighbours. The observation belongs to the class the most frequent amongst neighbours.



K-NEAREST NEIGHBOURS

When an observation comes in, calculate the distance to the k nearest neighbours. The observation belongs to the class the most frequent amongst neighbours.



Emotions?
oo

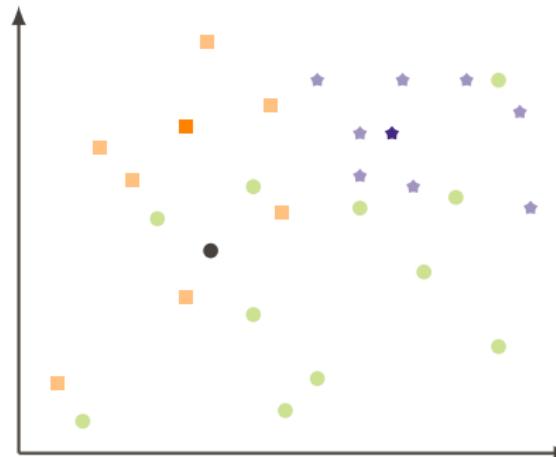
Classification
oooooooooooo●oooooooooooo

Emotion classifier
oooooooooooo

ML for social behaviours generation
oooooooooooo

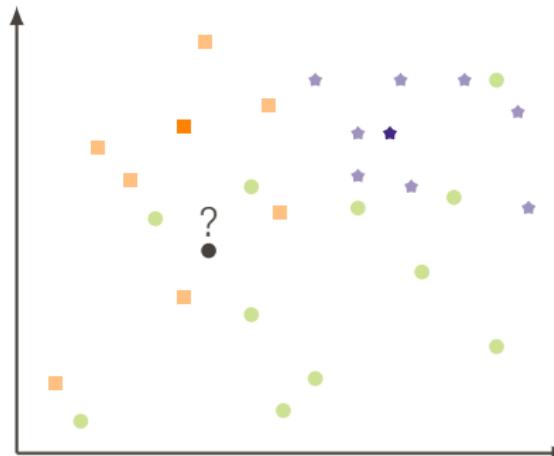
K-NEAREST NEIGHBOURS

When an observation comes in, calculate the distance to the k nearest neighbours. The observation belongs to the class the most frequent amongst neighbours.



K-NEAREST NEIGHBOURS

When an observation comes in, calculate the distance to the k nearest neighbours. The observation belongs to the class the most frequent amongst neighbours.



Emotions?
oo

Classification
oooooooo●ooooooooooooooo

Emotion classifier
ooooooooooooooo

ML for social behaviours generation
ooooooooooooooo

PYTHON IMPLEMENTATION OF K -NEAREST NEIGHBOURS

Complete example, with the data above:

```
from numpy import genfromtxt
from sklearn import neighbors

""" data.csv:
2.2,1.3,0 -> circles
3.2,2.3,0
...
2.3,3.2,1 -> squares
0.3,0.6,1
...
2.8,3.5,2 -> stars
3.2,2.6,2
...
"""

inputs = [ [3.5,3], [1.5,3], [1.8,1.9] ]
k = 3
knns = neighbors.KNeighborsClassifier(k)
knns.fit(data, categories)

predictions = knns.predict(inputs)
print(predictions)

>>> [2. 1. 1.]
```

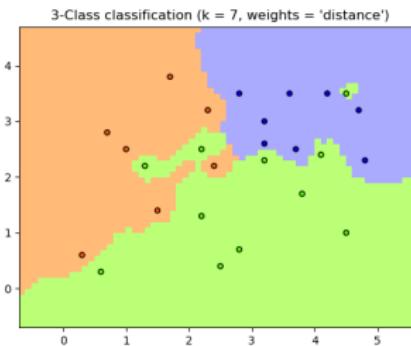
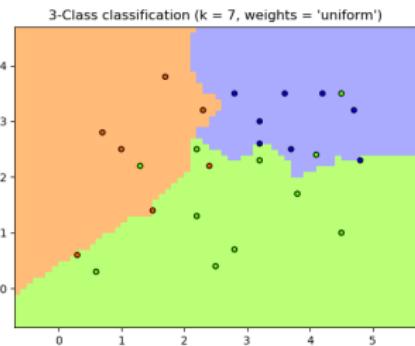
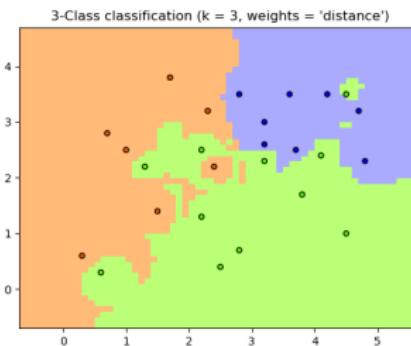
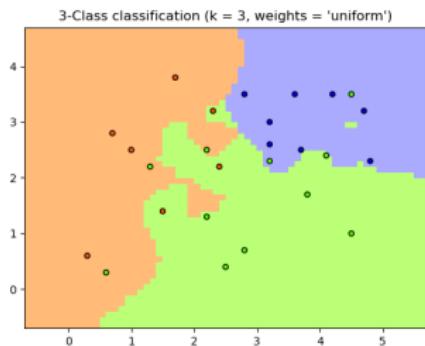
Emotions?
oo

Classification
oooooooo●ooooooooooooooo

Emotion classifier
ooooooooooooooo

ML for social behaviours generation
ooooooooooooooo

CHOICE OF K AND WEIGHTS



Emotions?
oo

Classification
oooooooo●ooooooooooooooo

Emotion classifier
ooooooooooooooo

ML for social behaviours generation
ooooooooooooooo

K-NEAREST NEIGHBOURS: SUMMARY

- *k*-nearest Neighbours usually does really well on a large number of classification problem, but can underperform near the classification boundary.

K-NEAREST NEIGHBOURS: SUMMARY

- *k*-nearest Neighbours usually does really well on a large number of classification problem, but can underperform near the classification boundary.
- might struggle with very large datasets, as it needs to calculate the distance to all training data every time you classify a new observation (workarounds exist, relying on approximate distances and optimisation of the algorithmic code)

K-NEAREST NEIGHBOURS: SUMMARY

- *k*-nearest Neighbours usually does really well on a large number of classification problem, but can underperform near the classification boundary.
- might struggle with very large datasets, as it needs to calculate the distance to all training data every time you classify a new observation (workarounds exist, relying on approximate distances and optimisation of the algorithmic code)
- commonly, the neighbours are weighted with the inverse of the distance $\frac{1}{d}$ to the observation (i.e. closer neighbours have stronger influence). With `sklearn`:
`neighbors.KNeighborsClassifier(k, weights='distance')`

Emotions?
oo

Classification
oooooooo●oooooooooooo

Emotion classifier
oooooooooooo

ML for social behaviours generation
oooooooooooo

SUPPORT VECTOR MACHINES

A very popular classification algorithm introduced by Vapnik in 1992.

Emotions?
oo

Classification
oooooooo●oooooooooooo

Emotion classifier
oooooooooooo

ML for social behaviours generation
oooooooooooo

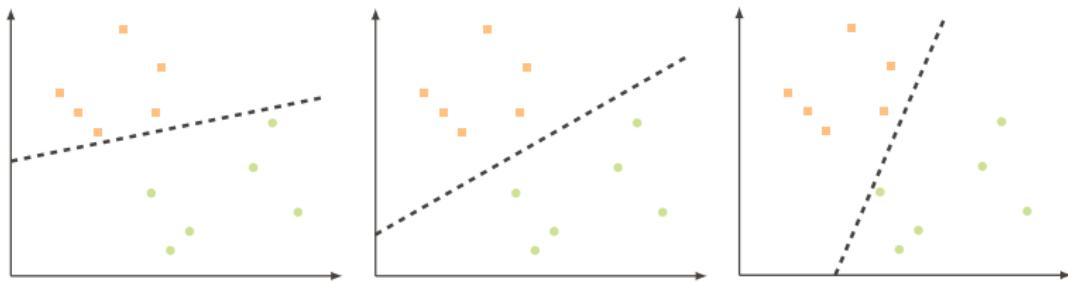
SUPPORT VECTOR MACHINES

A very popular classification algorithm introduced by Vapnik in 1992.

Often (but not always) provides very impressive classification performance on reasonably sized datasets.

SVM PRINCIPLE

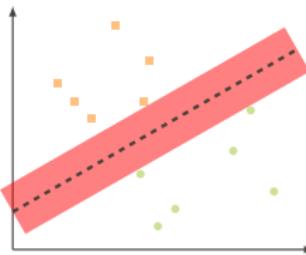
Three different classification lines. All are correct, but is there any reason why one is better than the others?



Intuitively, the line that is most distant to the training data is the best division.

SVM PRINCIPLE

Three different classification lines. All are correct, but is there any reason why one is better than the others?



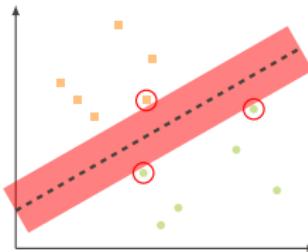
Intuitively, the line that is most distant to the training data is the best division.

The empty area near the division line is symmetric. It forms a bar in 2D space, a cylinder in 3D space, and a hyper-cylinder in n-D space.

SVM PRINCIPLE

The classifier in the middle is called the **maximum margin classifier**.

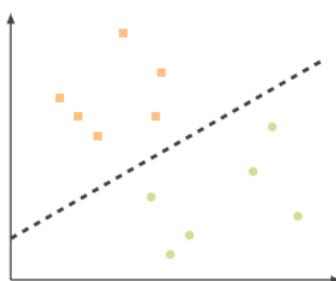
The data points nearest the classifier are called **support vectors**.



Two observations

- The margins should be as large as possible.
- The support vectors are the most useful datapoints because they are the ones that we might get wrong.

LINEAR CLASSIFIER



A linear classifier for two categories (**binary classifier**) can be written as $f(\mathbf{x}) = \mathbf{w} \cdot \mathbf{x} + b$.

- $\mathbf{w} \cdot \mathbf{x} = \sum_i w_i x_i$. This is called the scalar product or inner product. Can also be written as a matrix multiplication $\mathbf{w}^T \mathbf{x}$.
- \mathbf{w} is a **weight vector**, tilting the line
- \mathbf{x} is a data point (of dimension n)
- b is a **bias**, lifting the line up along the y -axis

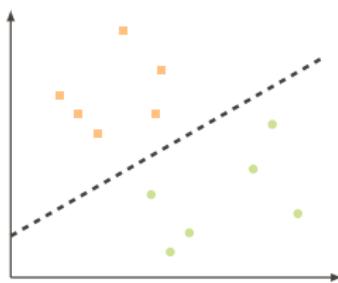
Emotions?
oo

Classification
oooooooooooo●oooooooooooo

Emotion classifier
oooooooooooooooo

ML for social behaviours generation
oooooooooooooooo

LINEAR CLASSIFIER

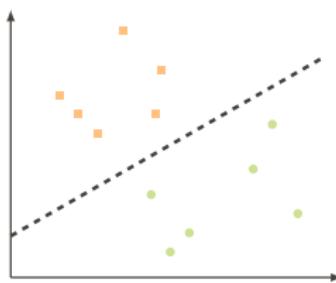


\mathbf{w} has to be calculated such that:

$$f(\mathbf{x}) < 0 \Rightarrow \mathbf{x} \in \{\blacksquare\}$$

$$f(\mathbf{x}) > 0 \Rightarrow \mathbf{x} \in \{\circ\}$$

LINEAR CLASSIFIER



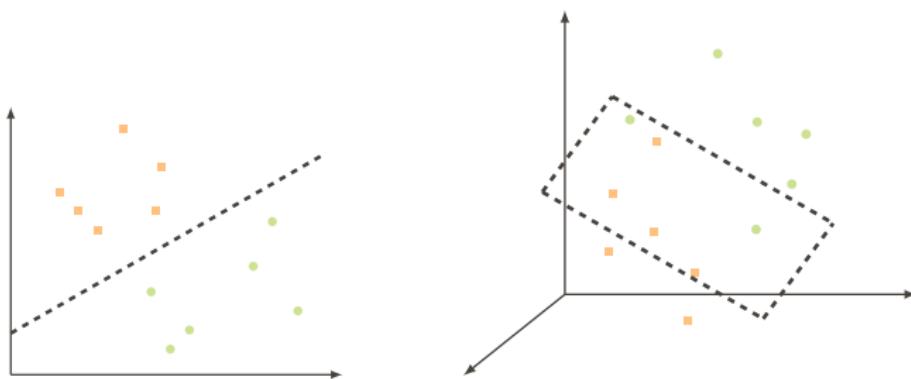
w has to be calculated such that:

for kNNs, we have to carry along the training data;
with a linear classifier,
once **w** is found, we can discard the training data

$$f(\mathbf{x}) < 0 \Rightarrow \mathbf{x} \in \{\blacksquare\}$$

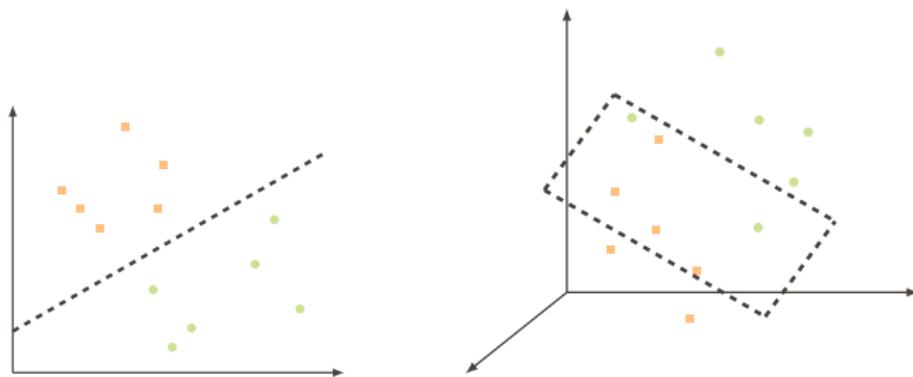
$$f(\mathbf{x}) > 0 \Rightarrow \mathbf{x} \in \{o\}$$

LINEAR CLASSIFIER



- In 2D, the **discriminant** is a line
- In 3D, the discriminant is a plane
- In n-D , the discriminant is a hyperplane

LINEAR CLASSIFIER



- In 2D, the **discriminant** is a line
- In 3D, the discriminant is a plane
- In n-D , the discriminant is a hyperplane

the 'dimensions' are the **features** of our inputs,
e.g. for a human, the size, age, skin colour, gender...

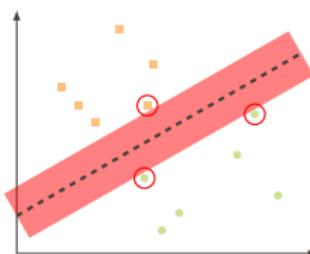
Emotions?
oo

Classification
oooooooooooo●oooooooooooo

Emotion classifier
oooooooooooooooo

ML for social behaviours generation
oooooooooooooooo

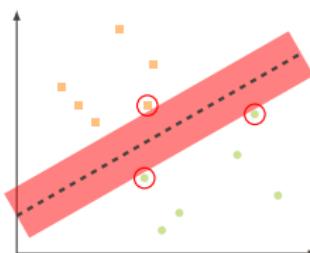
SUPPORT VECTOR CLASSIFIER



In a SVM, we want to **maximise the margin M** . We can rewrite the linear classifier.

If $\mathbf{w}^T \mathbf{x} + b \geq M$, then an observation belongs to ■, for $\mathbf{w}^T \mathbf{x} + b \leq -M$, it belongs to ○.

SUPPORT VECTOR CLASSIFIER



In a SVM, we want to **maximise the margin M** . We can rewrite the linear classifier.

If $\mathbf{w}^T \mathbf{x} + b \geq M$, then an observation belongs to ■, for $\mathbf{w}^T \mathbf{x} + b \leq -M$, it belongs to ○.

A point \mathbf{x}^\blacksquare that lies on the ■ class boundary line, i.e. $\mathbf{w}^T \mathbf{x}^\blacksquare + b = M$, is a **support vector**.

Emotions?
oo

Classification
oooooooooooo●oooooooooooo

Emotion classifier
oooooooooooo

ML for social behaviours generation
oooooooooooo

FINDING THE OPTIMAL CLASSIFIER

The problem is to find the **optimal** values for w and b .

The classifier needs to satisfy two conditions

- It needs to correctly classify the training data,
- and needs the margin from the classifier to be as large as possible.

FINDING THE OPTIMAL CLASSIFIER

The problem is to find the **optimal** values for \mathbf{w} and b .

The classifier needs to satisfy two conditions

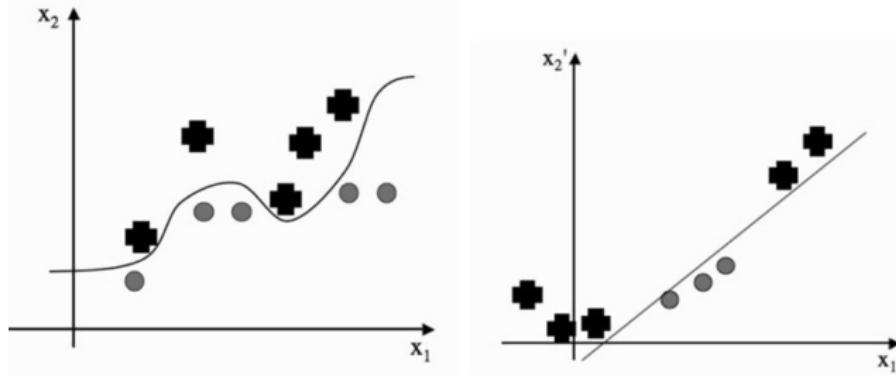
- It needs to correctly classify the training data,
- and needs the margin from the classifier to be as large as possible.

A **quadratic programming solver** is used to find the optimal values for \mathbf{w} and b .

Oxford lecture on SVM to learn more about the mathematical formulation and the *perceptron algorithm*.

TRANSFORMATION OF DATA

- Unfortunately, *all this still assumes that the data is linearly separable*. But what if it is not, and we are not prepared for a few misclassifications?
- The solution is to **transform** the data: move data points in the n-D space until the training data is linearly separable again



Emotions?
oo

Classification
oooooooooooo●ooooooo

Emotion classifier
oooooooooooo

ML for social behaviours generation
oooooooooooo

FEATURE MAPS

We can transform data points (= move them) or even add more dimensions using a function $\theta(\mathbf{x}_i)$ from input \mathbf{x}_i . θ is a **feature map**.

Example:



Emotions?
oo

Classification
oooooooooooo●ooooooo

Emotion classifier
oooooooooooo

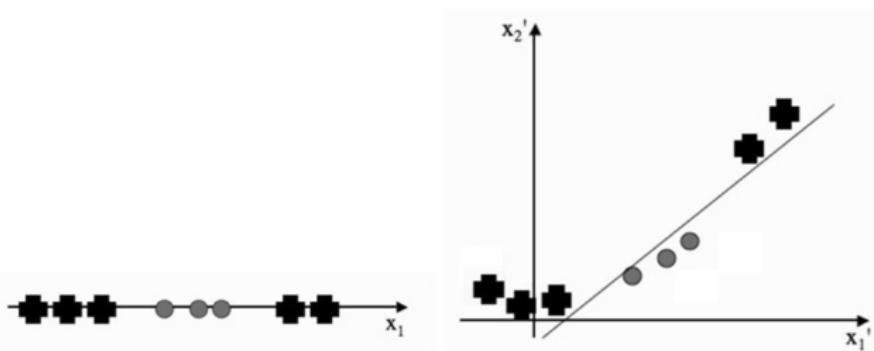
ML for social behaviours generation
oooooooooooo

FEATURE MAPS

We can transform data points (= move them) or even add more dimensions using a function $\theta(\mathbf{x}_i)$ from input \mathbf{x}_i . θ is a **feature map**.

Example:

$$\theta(\mathbf{x}_i) = [\mathbf{x}_i, \mathbf{x}_i^2]$$



→ now, linearly separable

Emotions?
oo

Classification
oooooooooooo●oooooooo

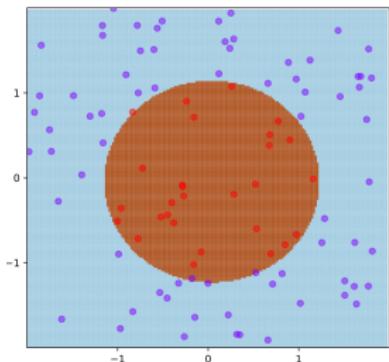
Emotion classifier
oooooooooooo

ML for social behaviours generation
oooooooooooo

FEATURE MAPS

We can transform data points (= move them) or even add more dimensions using a function $\theta(\mathbf{x}_i)$ from input \mathbf{x}_i . θ is a **feature map**.

Example:



Emotions?
oo

Classification
oooooooooooo●ooooooo

Emotion classifier
oooooooooooo

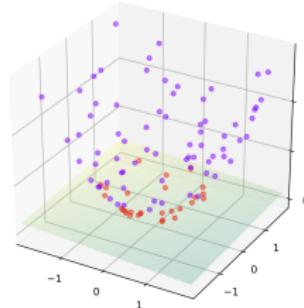
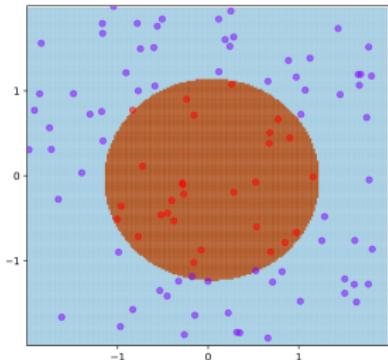
ML for social behaviours generation
oooooooooooo

FEATURE MAPS

We can transform data points (= move them) or even add more dimensions using a function $\theta(\mathbf{x}_i)$ from input \mathbf{x}_i . θ is a **feature map**.

Example:

$$\theta(<\mathbf{x}_i, \mathbf{y}_i>) = [\mathbf{x}_i, \mathbf{y}_i, \mathbf{x}_i^2 + \mathbf{y}_i^2]$$



→ now, linearly separable

Emotions?
oo

Classification
oooooooooooo●oooooooo

Emotion classifier
oooooooooooo

ML for social behaviours generation
oooooooooooo

FEATURE MAPS

If we know something about the structure of the data, then we might be able to identify feature maps that would be effective.

Often however we do not have such domain knowledge.

Emotions?
oo

Classification
oooooooooooo●ooooooo

Emotion classifier
oooooooooooo

ML for social behaviours generation
oooooooooooo

THE KERNEL TRICK

$$\theta : \mathbf{x} \rightarrow \theta(\mathbf{x}), \mathbb{R}^d \rightarrow \mathbb{R}^D$$

$$f(\mathbf{x}) = \mathbf{w}^T \theta(\mathbf{x}) + b$$

- o Simply map \mathbf{x} to $\theta(\mathbf{x})$ where data is separable
- o Solve for \mathbf{w} in high dimensional space \mathbb{R}^D
- o If $D \gg d$ then there are many more parameters to learn for \mathbf{w} . Can this be avoided?

Emotions?
oo

Classification
oooooooooooo●ooooooo

Emotion classifier
oooooooooooo

ML for social behaviours generation
oooooooooooo

THE KERNEL TRICK

$$\theta : \mathbf{x} \rightarrow \theta(\mathbf{x}), \mathbb{R}^d \rightarrow \mathbb{R}^D$$

$$f(\mathbf{x}) = \mathbf{w}^T \theta(\mathbf{x}) + b$$

It can be shown that you do not actually need to calculate explicitly $\theta(\mathbf{x})$. Instead, you only need $\theta(\mathbf{x}_i)^T \theta(\mathbf{x}_j)$ (with \mathbf{x}_i the training data) to compute \mathbf{w} .

We call the **kernel function** $K(\mathbf{x}_i, \mathbf{x}_j) = \theta(\mathbf{x}_i)^T \theta(\mathbf{x}_j)$. The SVM classifier can be learnt and applied with only K , and the complexity depends only on N (size of training data), not on D .

STANDARD KERNELS

- **Polynomials** up to some degree s in the elements x_k of the input vector (e.g. x_3^3 or x_1x_4). This can be written as:

$$K(x, y) = (1 + x^T y)^s$$

- **Sigmoid functions** of the x_k with parameters κ and δ , and kernel:

$$K(x, y) = \tanh(\kappa x^T y - \delta)$$

- **Radial basis function** expansions of the x_k with parameter σ and kernel:

$$K(x, y) = \exp\left(-\frac{(x - y)^2}{2\sigma^2}\right)$$

(a Gaussian kernel)

Emotions?
oo

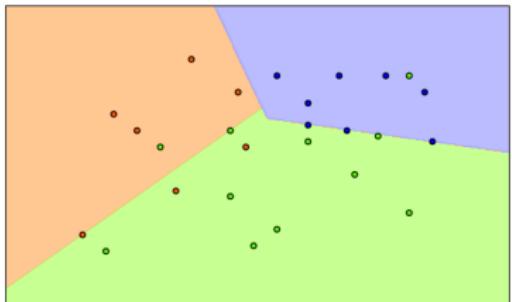
Classification
oooooooooooooooooooo●ooooo

Emotion classifier
ooooooooooooooo

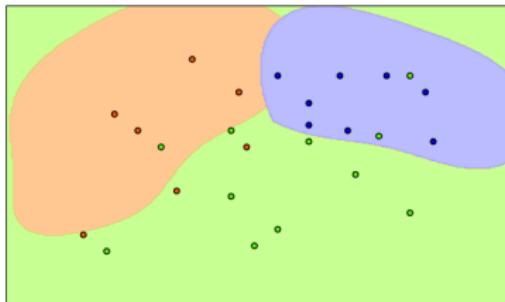
ML for social behaviours generation
ooooooooooooooo

COMPARISON OF KERNEL FUNCTIONS

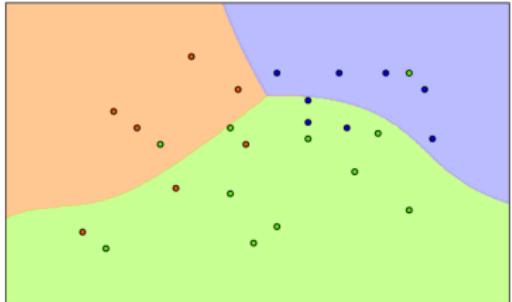
SVM with linear kernel



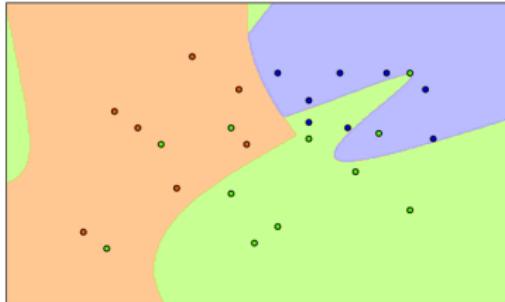
SVM with RBF kernel



SVM with polynomial (degree 3) kernel



SVM with polynomial (degree 6) kernel



Emotions?
oo

Classification
oooooooooooooooooooo●oooo

Emotion classifier
oooooooooooo

ML for social behaviours generation
oooooooooooo

PYTHON IMPLEMENTATION OF SUPPORT VECTOR MACHINE

```
from sklearn import svm

C = 1.0 # SVM regularization parameter
          # +- 'how acceptable are misclassification'
clf = svm.SVC(kernel='linear', C=C) # kernel='poly', 'rbf', 'sigmoid'...

clf.fit(data, categories)

predictions = clf.predict(inputs)
```

Emotions?
oo

Classification
oooooooooooooooooooo●oooo

Emotion classifier
oooooooooooo

ML for social behaviours generation
oooooooooooo

PYTHON IMPLEMENTATION OF SUPPORT VECTOR MACHINE

Complete example, with our data:

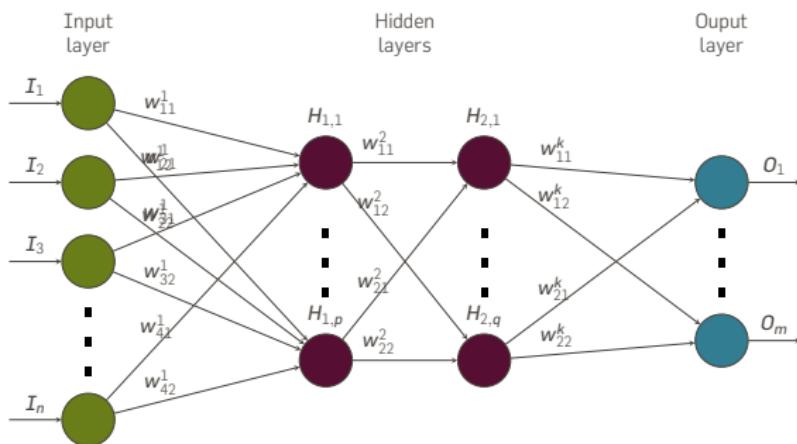
```
from numpy import genfromtxt
from sklearn import svm

""" data.csv:
2.2,1.3,0 -> circles
3.2,2.3,0
...
2.3,3.2,1 -> squares
0.3,0.6,1
...
2.8,3.5,2 -> stars
3.2,2.6,2
...
"""
inputs = [ [3.5,3], [1.5,3], [1.8,1.9] ]
clf = svm.SVC(kernel='rbf',
                gamma = 0.7,
                C=1.0)
clf.fit(data, categories)

predictions = clf.predict(inputs)
print(predictions)

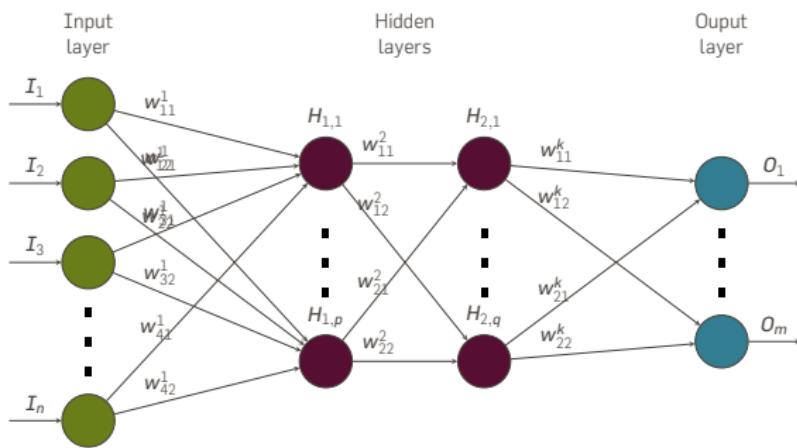
>>> [2. 1. 0.]
```

MULTI-LAYER PERCEPTRONS (MLP)



- *fully connected feed-forward* artificial neural network
- one *input layer*, n ($n >= 1$) *hidden layers*, one *output layer*
- between each layer, a non-linear *activation function* (often sigmoid or ReLU)

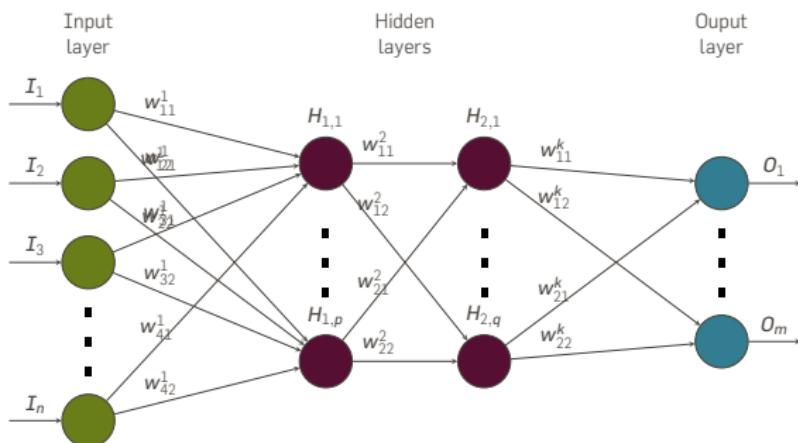
MULTI-LAYER PERCEPTRONS (MLP)



In our case:

- 2 coordinates: size n of the input layer: 2
- 3 categories: size m of the output layer: 1 (or 3 if using *one-hot encoding* for categories)
- nb and size of hidden layers? **hyper-parameters**. Need to be determined as well! (\rightarrow validation set)

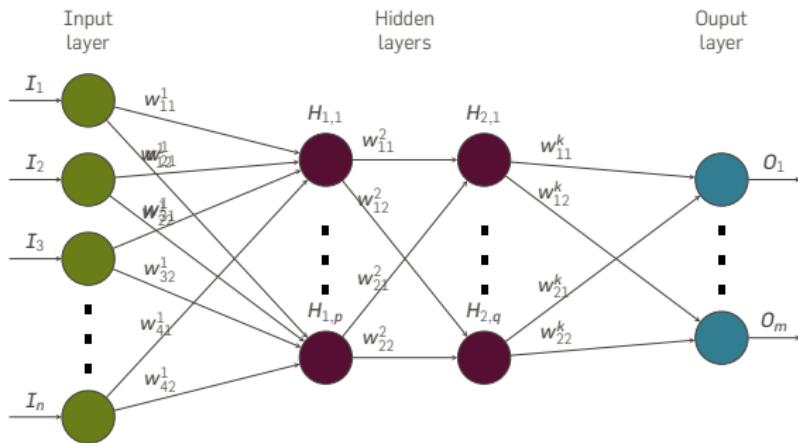
MULTI-LAYER PERCEPTRONS (MLP)



For each layer k : $\mathbf{x}^k = f(\mathbf{W}^T \mathbf{x}^{k-1} + \mathbf{b})$

- \mathbf{x}^{k-1} the input vector of this layer
- \mathbf{x}^k the output vector
- \mathbf{W} the set of parameters or weights
- \mathbf{b} the bias
- f the activation function (eg ReLU, sigmoid)

MULTI-LAYER PERCEPTRONS (MLP)



Training \Rightarrow finding weights w_{ij}^k so that the outputs values match 'what we expect'.

Main algorithm: *backpropagation* using *gradient descent*

Emotions?
oo

Classification
oooooooooooooooooooo●oo

Emotion classifier
oooooooooooooo

ML for social behaviours generation
oooooooooooooo

PYTHON IMPLEMENTATION OF MULTI-LAYER PERCEPTRONS

```
from numpy import genfromtxt

from sklearn.neural_network import MLPClassifier

csv = genfromtxt('data.csv', delimiter=',')
data = csv[:, :2]
categories = csv[:, 2]

inputs = [ [3.5, 3], [1.5, 3], [1.8, 1.9] ]

clf = MLPClassifier(hidden_layer_sizes=(4, 4),
                     activation = "relu")
clf.fit(data, categories)

predictions = clf.predict(inputs)
print(predictions)

>>> [2. 2. 0.]
```

Emotions?
oo

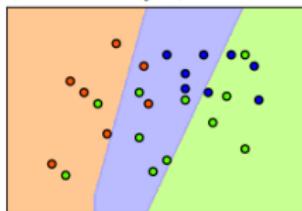
Classification
oooooooooooooooooooo●●

Emotion classifier
oooooooooooo

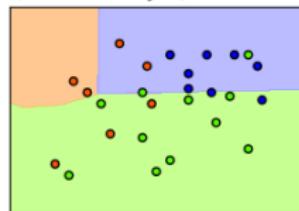
ML for social behaviours generation
oooooooooooo

MLP HYPER-PARAMETERS

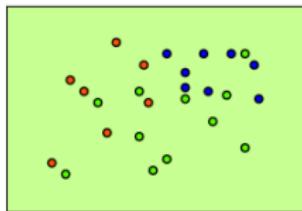
MLP, 1 hidden layer, 4 hidden units



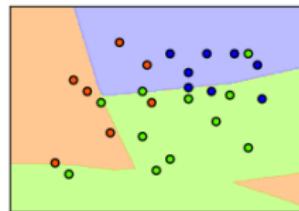
MLP, 2 hidden layer, 4 hidden units



MLP, 3 hidden layer, 4 hidden units



MLP, 1 hidden layer, 8 hidden units



Emotions?
oo

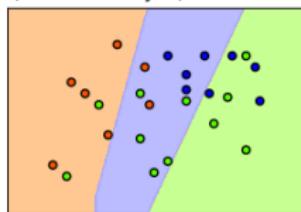
Classification
oooooooooooooooooooo●●

Emotion classifier
oooooooooooo

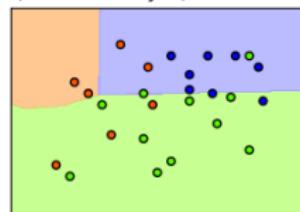
ML for social behaviours generation
oooooooooooo

MLP HYPER-PARAMETERS

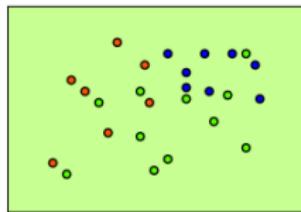
MLP, 1 hidden layer, 4 hidden units



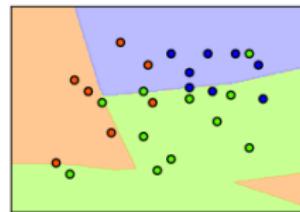
MLP, 2 hidden layer, 4 hidden units



MLP, 3 hidden layer, 4 hidden units



MLP, 1 hidden layer, 8 hidden units



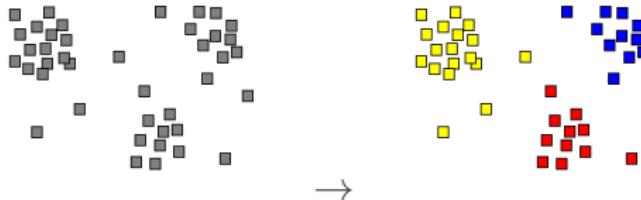
MLP does not perform well on this example: not enough data to converge!

CLUSTERING VS CLASSIFICATION

Clustering is sometimes confused with **classification**. (often because some algorithms have similar names, e.g. *k-means clustering* and *k-nearest neighbours*).

Clustering starts from unlabelled data points and tries to find k clusters in the data → **unsupervised learning**.

→ Used to find patterns in the data.



Classification starts from training data (→ **supervised learning**), and attempts to correctly classify an unknown observation.

LET'S BUILD OUR OWN EMOTION
CLASSIFIER FROM SCRATCH

Emotions?
oo

Classification
oooooooooooooooooooo

Emotion classifier
o●oooooooooooo

ML for social behaviours generation
oooooooooooo

SHOPPING LIST

What do we need?

Emotions?
oo

Classification
oooooooooooooooooooo

Emotion classifier
o●oooooooooooo

ML for social behaviours generation
oooooooooooo

SHOPPING LIST

What do we need?

- a dataset of labelled faces

Emotions?
oo

Classification
oooooooooooooooooooo

Emotion classifier
o●oooooooooooo

ML for social behaviours generation
oooooooooooo

SHOPPING LIST

What do we need?

- a dataset of labelled faces
- some pre-processing to normalise the faces

Emotions?
oo

Classification
oooooooooooooooooooo

Emotion classifier
o●oooooooooooo

ML for social behaviours generation
oooooooooooo

SHOPPING LIST

What do we need?

- a dataset of labelled faces
- some pre-processing to normalise the faces
- features to extract

Emotions?
oo

Classification
oooooooooooooooooooo

Emotion classifier
o●oooooooooooo

ML for social behaviours generation
oooooooooooo

SHOPPING LIST

What do we need?

- a dataset of labelled faces
- some pre-processing to normalise the faces
- features to extract
- a classifier



human face happiness



16



All

Images

Videos

Shopping

News

More

Settings

Tools

View saved

SafeSearch ▾

smile

cartoon

male

boy

grandpa

black and white

real

happiness

joy

sad

gratitude

beauty

emotional

smiley

hawa >







Emotions?

oo

Classification

oooooooooooooooooooo

Emotion classifier

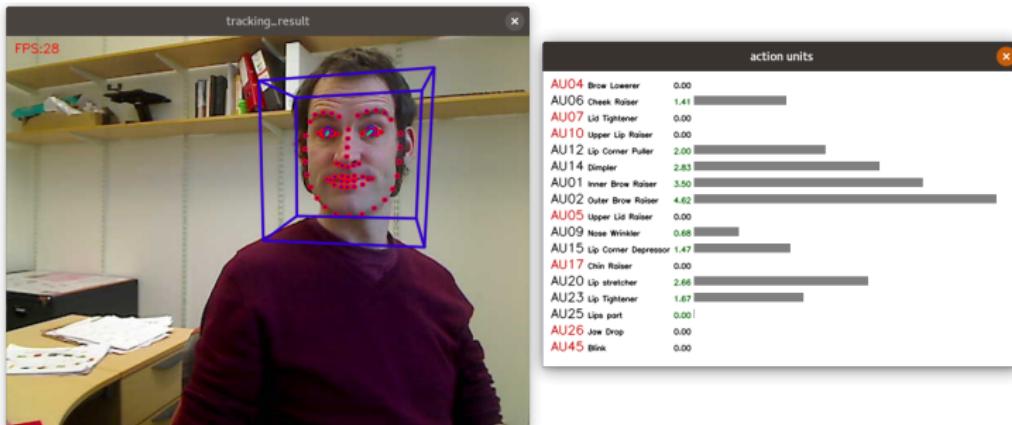
oooo●oooo

ML for social behaviours generation

oooooooooooo

ACTION UNITS

Our features: **action units.**



github.com/TadasBaltrusaitis/OpenFace

| | A | B | C | D | E | F | G | H | I | J | K | L | M | N | O | P | Q | R | S | |
|-----|---------|------|------|------|------|------|------|------|------|------|------|------|------|------|------|------|------|------|--|---|
| 1 | emotion | AU01 | AU02 | AU04 | AU05 | AU06 | AU07 | AU09 | AU10 | AU12 | AU14 | AU15 | AU17 | AU20 | AU23 | AU25 | AU26 | AU45 | face | |
| 2 | 0 | 1.07 | 0 | 0.05 | 2.95 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1.36 | 0 | 0 | fear/processed/aligned/fear-041_aligned.bmp | |
| 3 | 0 | 1.57 | 0 | 0.63 | 0 | 1.02 | 1.39 | 0 | 0.45 | 0 | 0 | 0 | 0 | 0 | 0 | 0.45 | 1.35 | 0 | fear/processed/aligned/fear-052_aligned.bmp | |
| 4 | 0 | 1.38 | 0 | 1.17 | 0 | 1.36 | 0 | 1.85 | 1.21 | 0 | 0 | 0 | 0 | 0.05 | 0 | 0 | 0 | 0 | fear/processed/aligned/fear-035_aligned.bmp | |
| 5 | 0 | 0 | 0.97 | 0.87 | 2.48 | 0 | 0 | 0 | 0 | 0 | 0 | 0.31 | 0 | 0.71 | 0.32 | 1.41 | 0.24 | 0 | fear/processed/aligned/fear-018_aligned.bmp | |
| 6 | 0 | 3.79 | 2.52 | 0 | 4.61 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1.53 | 0.67 | 0 | fear/processed/aligned/fear-043_aligned.bmp | |
| 7 | 0 | 2.55 | 0 | 0 | 1.2 | 1.58 | 0 | 0 | 1.27 | 0.64 | 0 | 0.93 | 1.14 | 0 | 0 | 0 | 0 | 0 | fear/processed/aligned/fear-044_aligned.bmp | |
| 8 | 0 | 1.93 | 2.72 | 0 | 3.86 | 0 | 0 | 0 | 0 | 0 | 0 | 1.38 | 0.67 | 1.23 | 0 | 0.15 | 0.83 | 0 | fear/processed/aligned/fear-006_aligned.bmp | |
| 75 | 1 | 0 | 1 | 0 | 0 | 1.79 | 3.94 | 0 | 2.96 | 2.62 | 0 | 0.79 | 0 | 0 | 0 | 0.36 | 0 | 0 | 1.08 happiness/processed/aligned/happiness-098_aligned.bmp | |
| 76 | 1 | 0 | 0 | 0 | 1.78 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 2.14 | 1.55 | 0 | happiness/processed/aligned/happiness-093_aligned.bmp | |
| 77 | 1 | 0 | 0 | 0 | 0 | 2.17 | 2.74 | 1.8 | 1.83 | 1.91 | 1.23 | 0 | 0 | 0.61 | 0 | 3.24 | 0 | 0 | 0.04 happiness/processed/aligned/happiness-095_aligned.bmp | |
| 78 | 1 | 0 | 0 | 2.91 | 0 | 1.38 | 0 | 2.71 | 2.11 | 1.04 | 1.38 | 0.15 | 0.71 | 1.4 | 0.23 | 0.56 | 0 | 0 | 0.65 happiness/processed/aligned/happiness-046_aligned.bmp | |
| 79 | 1 | 0.51 | 0 | 0 | 0 | 1.23 | 2.38 | 0.82 | 0.33 | 2.21 | 0.99 | 0.03 | 0 | 0 | 0 | 1.5 | 0.39 | 0 | 0.07 happiness/processed/aligned/happiness-025_aligned.bmp | |
| 80 | 1 | 0.51 | 0.76 | 0 | 0 | 1.56 | 0 | 0 | 1.23 | 2.56 | 1.12 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 happiness/processed/aligned/happiness-024_aligned.bmp |
| 81 | 1 | 0 | 0 | 0 | 0 | 3.22 | 3.45 | 0 | 2.84 | 3.37 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 happiness/processed/aligned/happiness-075_aligned.bmp |
| 196 | 2 | 0 | 0.1 | 3.45 | 1.14 | 0 | 0 | 1.75 | 0.65 | 0 | 0 | 0.11 | 1.04 | 2.05 | 0 | 0.67 | 0 | 0 | 0.12 anger/processed/aligned/anger-051_aligned.bmp | |
| 197 | 2 | 0 | 0.92 | 0 | 0 | 0.43 | 1.79 | 0 | 0.24 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 anger/processed/aligned/anger-027_aligned.bmp |
| 198 | 2 | 1.27 | 0 | 2.58 | 0 | 0 | 0 | 0 | 0.39 | 0 | 0 | 1.35 | 2.39 | 1.28 | 0.28 | 0 | 0 | 0 | 0 | 0 anger/processed/aligned/anger-057_aligned.bmp |
| 199 | 2 | 4.5 | 1.72 | 0 | 1.24 | 1.54 | 0 | 0 | 1.63 | 0 | 0 | 0.22 | 0 | 0 | 0 | 2.29 | 0 | 0 | 0 | 0 anger/processed/aligned/anger-047_aligned.bmp |
| 200 | 2 | 0 | 0 | 0 | 0.29 | 0 | 0 | 0.56 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1.5 | 0 | 0 | 0 | 0 anger/processed/aligned/anger-008_aligned.bmp |
| 201 | 2 | 0 | 0 | 0 | 0.93 | 0.86 | 0 | 0.8 | 1.35 | 0.85 | 0 | 0 | 0 | 0 | 0 | 0.5 | 0 | 0 | 0 | 0 anger/processed/aligned/anger-031_aligned.bmp |
| 202 | 2 | 2.73 | 0 | 0 | 1.11 | 1.55 | 0.32 | 3.07 | 1.75 | 1.21 | 1.34 | 1.15 | 0 | 0 | 0.03 | 2.07 | 0 | 0 | 0 | 0 anger/processed/aligned/anger-019_aligned.bmp |
| 203 | 2 | 0 | 1.1 | 3.59 | 0 | 2.33 | 3.16 | 2.56 | 3.38 | 0 | 1.34 | 0 | 0 | 1.92 | 0 | 3.03 | 0 | 0 | 0 | 0 anger/processed/aligned/anger-016_aligned.bmp |
| 280 | 3 | 0 | 0 | 0 | 0.33 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0.59 | 0 | 0 | 0 | 0 | 0 surprise/processed/aligned/surprise-077_aligned.bmp |
| 281 | 3 | 0.34 | 0 | 0 | 1.38 | 1.24 | 0 | 0 | 1.48 | 2.05 | 0 | 0.57 | 0 | 0 | 0.12 | 3.27 | 0 | 0 | 0 | 0 surprise/processed/aligned/surprise-079_aligned.bmp |
| 282 | 3 | 1.44 | 0 | 0 | 0.59 | 0 | 0 | 0 | 1.66 | 0.76 | 0 | 0.49 | 0 | 0 | 0.08 | 1.31 | 0 | 0 | 0 | 0 surprise/processed/aligned/surprise-040_aligned.bmp |
| 283 | 3 | 0.24 | 0 | 0.9 | 1.43 | 0 | 0.89 | 0 | 1.23 | 0 | 0 | 1.29 | 0 | 0.56 | 0.11 | 0 | 0 | 0 | 0 surprise/processed/aligned/surprise-004_aligned.bmp | |
| 284 | 3 | 0.95 | 0 | 0 | 0 | 0 | 0 | 1.04 | 0 | 0 | 0.13 | 0 | 1.07 | 2.53 | 0.36 | 0 | 0 | 0 | 0 surprise/processed/aligned/surprise-062_aligned.bmp | |
| 285 | 3 | 0.69 | 0.91 | 0 | 1.87 | 0 | 0 | 0 | 1.34 | 0 | 0 | 1.09 | 0 | 0 | 0 | 2.94 | 1.45 | 0 | 0 | 0 surprise/processed/aligned/surprise-005_aligned.bmp |
| 286 | 3 | 1.19 | 0.68 | 0 | 2.03 | 0.33 | 0 | 0 | 0.58 | 1.73 | 1.65 | 0.22 | 0 | 0 | 0 | 1.92 | 0.7 | 0 | 0 | 0 surprise/processed/aligned/surprise-029_aligned.bmp |
| 287 | 3 | 3.88 | 2.83 | 0 | 4.25 | 0 | 0 | 0 | 0.52 | 1.14 | 0 | 1.09 | 0 | 0 | 0.1 | 1.49 | 1.64 | 0 | 0 | 0 surprise/processed/aligned/surprise-017_aligned.bmp |
| 288 | 3 | 0 | 0 | 0 | 1.48 | 1.46 | 0.79 | 0 | 2.65 | 1.39 | 0 | 1.08 | 0.4 | 0 | 0.89 | 0.11 | 0 | 0 | 0 | 0 surprise/processed/aligned/surprise-080_aligned.bmp |
| 289 | 3 | 0 | 0 | 0 | 0.79 | 1.62 | 0 | 2.49 | 0 | 0 | 0 | 0 | 0 | 0 | 0.01 | 1.62 | 1.87 | 0 | 0 | 0 surprise/processed/aligned/surprise-060_aligned.bmp |
| 290 | 3 | 1.87 | 0.79 | 0 | 2.75 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0.56 | 1.76 | 0.58 | 0 | 0 | 0 surprise/processed/aligned/surprise-002_aligned.bmp |
| 368 | 4 | 0 | 0 | 0 | 1.58 | 0.84 | 0 | 0 | 0 | 0 | 0.59 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 sadness/processed/aligned/sadness-025_aligned.bmp |

OUR TINY 'EMOTIONS' DATASET

- A couple of **Google queries**
- a bit of manual filtering
- 558 faces, between 70 and 130 per emotion
- split into 2 datasets: a **training** dataset (50 face per emotion) and a **testing** dataset
- for each face, OpenFace pre-processing to:
 - extract facial landmarks
 - normalise and mask out the faces
 - extract action units

The dataset is available for you to play with on the lecture's [GitHub repo](#).

Emotions?
oo

Classification
oooooooooooooooooooo

Emotion classifier
oooooooo●oooo

ML for social behaviours generation
oooooooooooo

CLASSIFICATION?

We want to train a classifier with our training dataset, and then check whether the emotions in the test faces can be properly recognised.

What are our options here?

Emotions?
oo

Classification
oooooooooooooooooooo

Emotion classifier
oooooooo●oooo

ML for social behaviours generation
oooooooooooo

CLASSIFICATION?

We want to train a classifier with our training dataset, and then check whether the emotions in the test faces can be properly recognised.

What are our options here?

1. kNN/SVM/MLP on the action units directly?

Emotions?
oo

Classification
oooooooooooooooooooo

Emotion classifier
oooooooo●oooo

ML for social behaviours generation
oooooooooooo

CLASSIFICATION?

We want to train a classifier with our training dataset, and then check whether the emotions in the test faces can be properly recognised.

What are our options here?

1. kNN/SVM/MLP on the action units directly?
2. PCA, then kNN/SVM/MLP on the action units?

Emotions?
oo

Classification
oooooooooooooooooooo

Emotion classifier
oooooooo●oooo

ML for social behaviours generation
oooooooooooo

CLASSIFICATION?

We want to train a classifier with our training dataset, and then check whether the emotions in the test faces can be properly recognised.

What are our options here?

1. kNN/SVM/MLP on the action units directly?
2. PCA, then kNN/SVM/MLP on the action units?
3. LDA, then kNN/SVM/MLP on the action units?

Emotions?
oo

Classification
oooooooooooooooooooo

Emotion classifier
oooooooo●oooo

ML for social behaviours generation
oooooooooooo

CLASSIFICATION?

We want to train a classifier with our training dataset, and then check whether the emotions in the test faces can be properly recognised.

What are our options here?

1. kNN/SVM/MLP on the action units directly?
2. PCA, then kNN/SVM/MLP on the action units?
3. LDA, then kNN/SVM/MLP on the action units?
4. Working with the (Eigen-)faces directly (eg, the pixels)?

Let's open a notebook!

(link to the complete notebook)

Emotions?
oo

Classification
oooooooooooooooooooo

Emotion classifier
oooooooooooo●ooo

ML for social behaviours generation
ooooooooooooooo

MACHINE LEARNING METRICS

| | | Predicted class | |
|--------------|-----------|-----------------|----------------|
| | | Class=yes | Class=no |
| Actual class | Class=yes | True positive | False negative |
| | Class=no | False positive | True negative |

Accuracy: $\frac{\text{correct predictions}}{\text{total observations}} = \frac{TP+TN}{TP+FP+FN+TN}$

Precision: $\frac{TP}{TP+FP}$

→ Did we assign too many emotions to a face?

Recall (eg sensitivity): $\frac{TP}{TP+FN}$

→ Did we 'miss' many emotions that were present?

F1 score: $2 \times \frac{\text{recall} \times \text{precision}}{\text{recall} + \text{precision}}$

Emotions?
oo

Classification
oooooooooooooooooooo

Emotion classifier
oooooooooooo●o

ML for social behaviours generation
oooooooooooo

HOW TO IMPROVE IT FURTHER?

Ideas?

Emotions?
oo

Classification
oooooooooooooooooooo

Emotion classifier
oooooooooooo●o

ML for social behaviours generation
oooooooooooo

HOW TO IMPROVE IT FURTHER?

Ideas?

- better dataset (wider range of facial expressions)
- larger dataset
- use the raw images directly? use the facial landmark?
- different classifier → neural network (deep or not)?



Natural emotions are much harder

ML FOR SOCIAL BEHAVIOURS GENERATION

LEARNING FROM DEMONSTRATION

- acquire new skills by learning to imitate an expert
- most often used for manipulation
- 3 main approaches:
kinesthetic teaching, teleoperation, visual observation
- Common algorithm:
reinforcement learning



Video source: Kormushev et al. 2010

a Kinesthetic teaching



b Teleoperation



c Passive observation



Emotions?
oo

Classification
oooooooooooooooooooo

Emotion classifier
oooooooooooo

ML for social behaviours generation
oo●oooooooooooo

LEARNING SEQUENCES

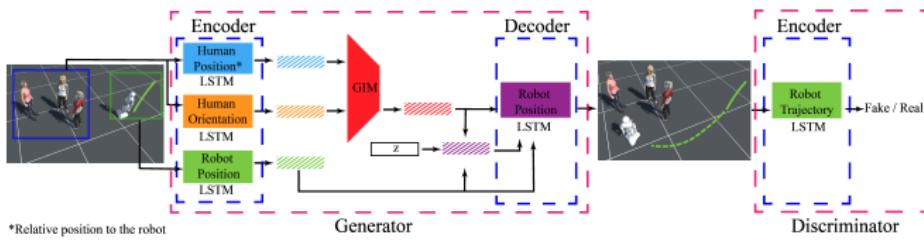
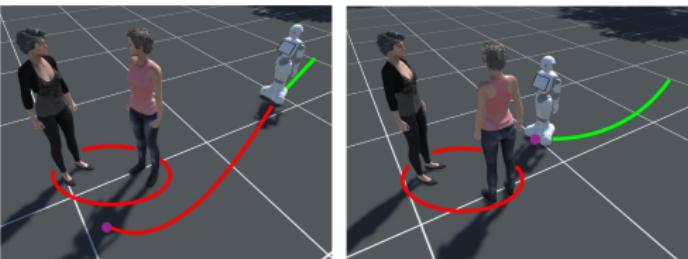


- 2 arms, 12 DoF
- Inputs: on-board 112×112 px camera, joint state
- Training from 40 teleoperated demonstrations of ≈ 70 s

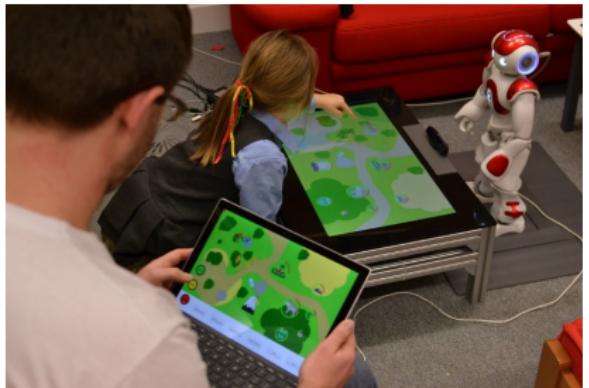
Not only learning poses, but **sequences as well**; *Time-Delay Neural Network* (TDNN) to learn to predict the next step.

GENERATIVE ADVERSARIAL NETS FOR BEHAVIOUR GENERATION

Some of the most exciting recent work in using ML for robot behaviour generation involve **Generative Adversarial Networks** (GANs):



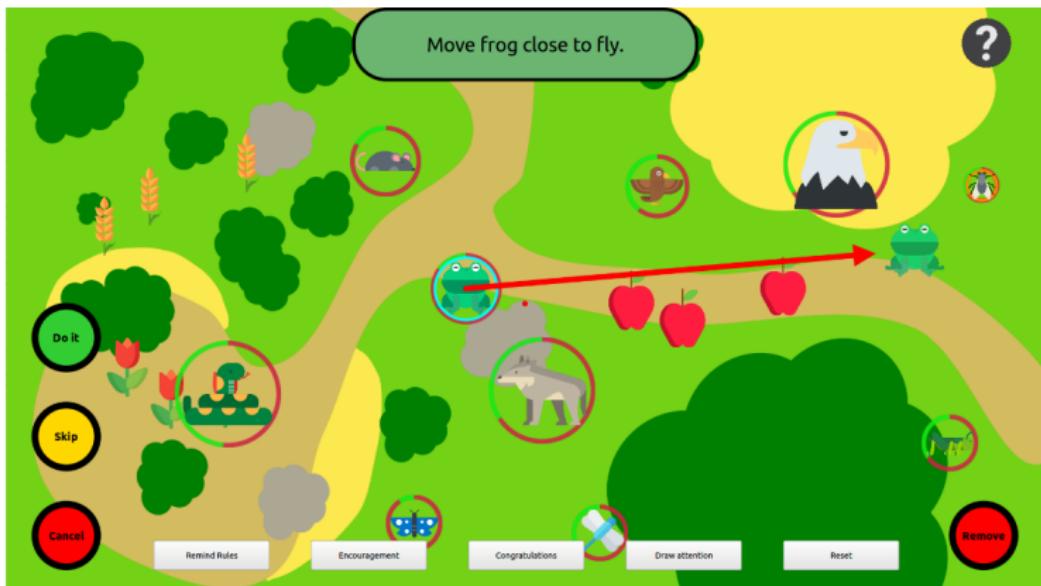
REINFORCEMENT LEARNING APPLIED TO SOCIAL ROBOTICS



The children plays a game about food chains; the robot learns to guide them (*task-specific action policy*) and encourage them (*social action policy*)

Complex problem: $|state| = 210$ $|action_space| = 655$

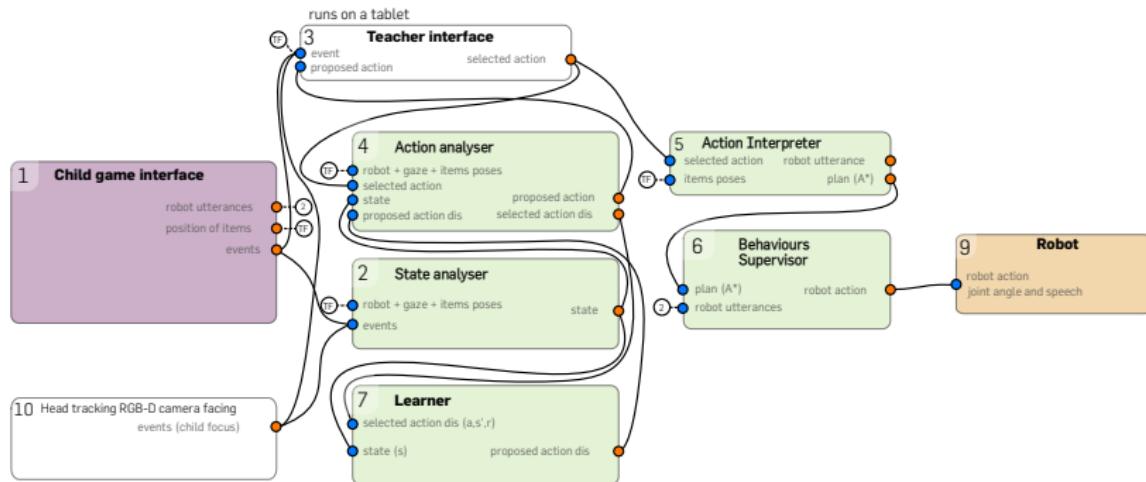
TEACHER'S INTERFACE



The robot's teacher (an end-user: might be the actual child's teacher) has a tablet interface that mirrors the child one, and adds robot's teleoperation and rewards.

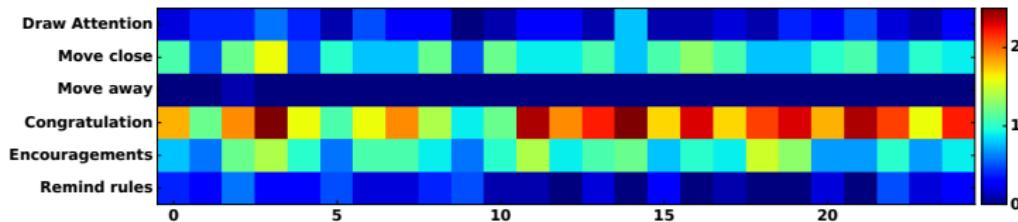


SOFTWARE ARCHITECTURE



LEARNT ROBOT'S BEHAVIOUR

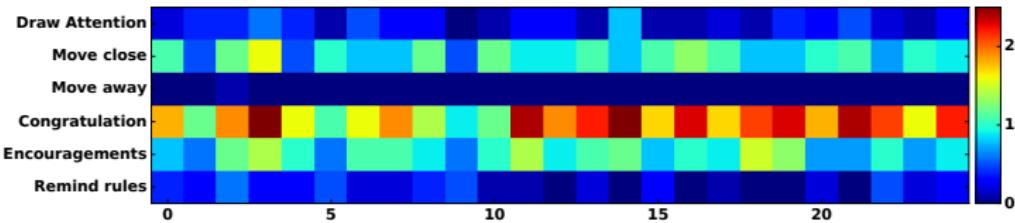
Distribution of actions for the 25 children participants:
Supervised



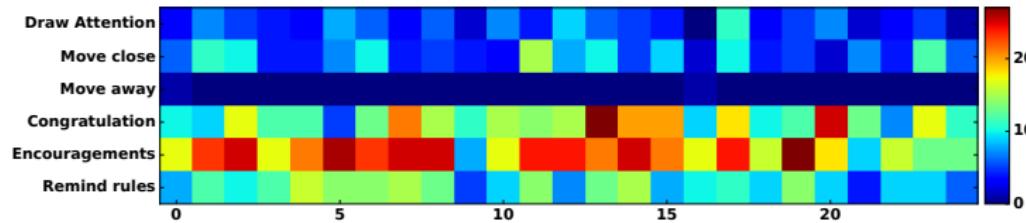
the robot personalises its action policies to the child's behaviour.

LEARNT ROBOT'S BEHAVIOUR

Distribution of actions for the 25 children participants:
Supervised



Autonomous

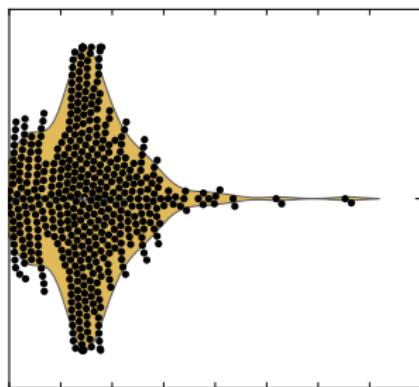


the robot personalises its action policies to the child's behaviour.

LEARNT ROBOT'S BEHAVIOUR

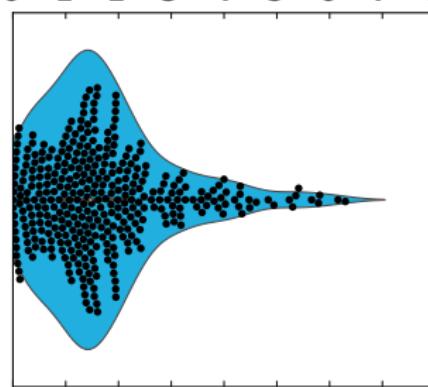
Time between a child's successful action and a praise:

0 1 2 3 4 5 6 7 8



supervised

0 1 2 3 4 5 6 7 8



autonomous

→ the robot has also learnt an appropriate 'social timing'.

Emotions?
oo

Classification
oooooooooooooooooooo

Emotion classifier
oooooooooooo

ML for social behaviours generation
oooooooo●oooo

COUCH-TO-5KM: EXPERT-IN-THE-LOOP MACHINE LEARNING



Emotions?
oo

Classification
oooooooooooooooooooo

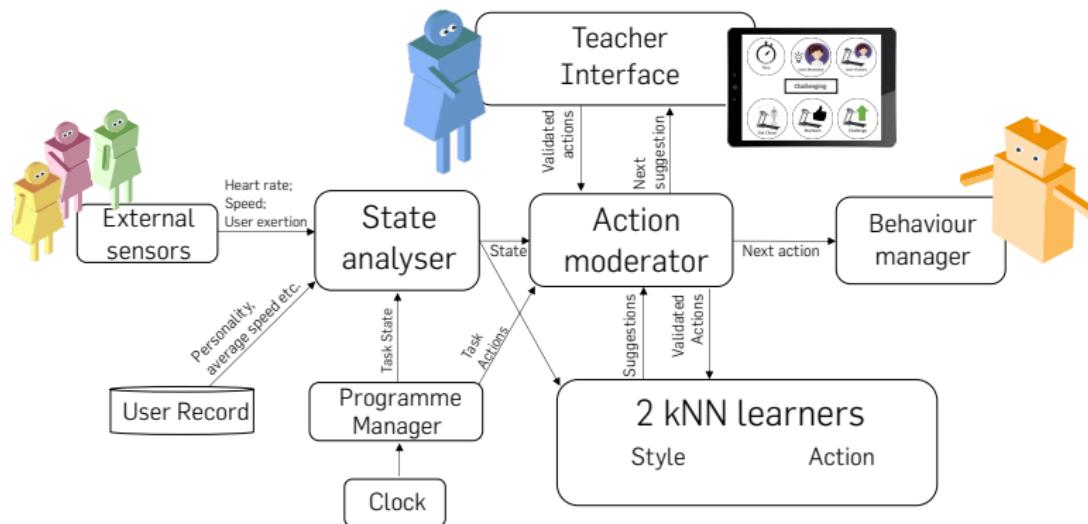
Emotion classifier
oooooooooooo

ML for social behaviours generation
oooooooo●ooo

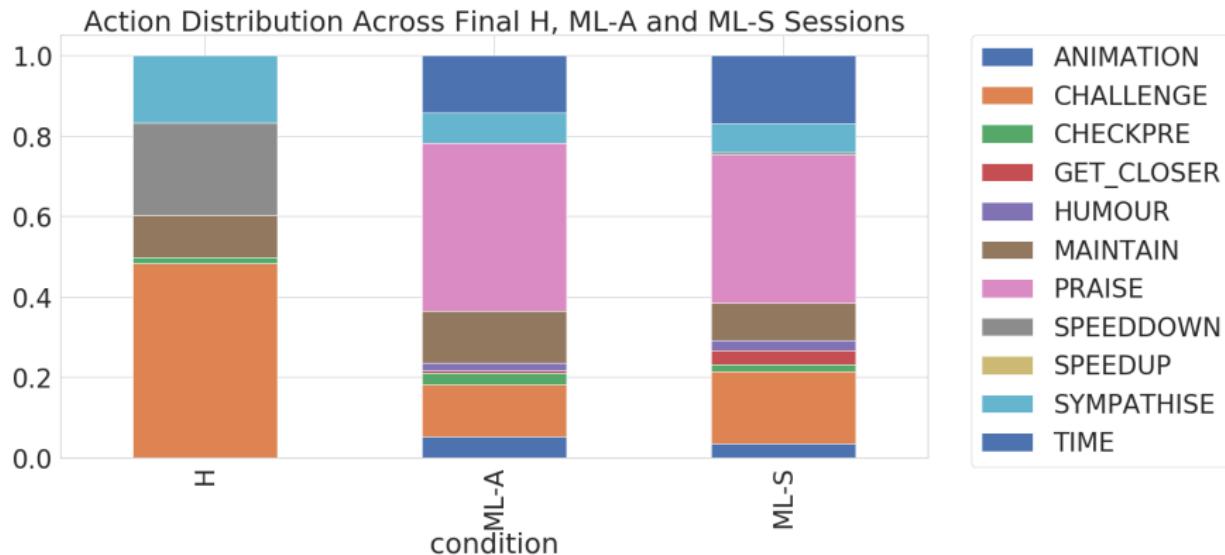
COUCH TO 5KM STUDY

- 9 participants
- 3 months; 27 one-hour sessions per participants
- 20 input features; 11 actions (task-specific or social)
- human-in-loop design and machine learning
- robot evolving from full teleoperation to full task and social autonomy

SOFTWARE ARCHITECTURE

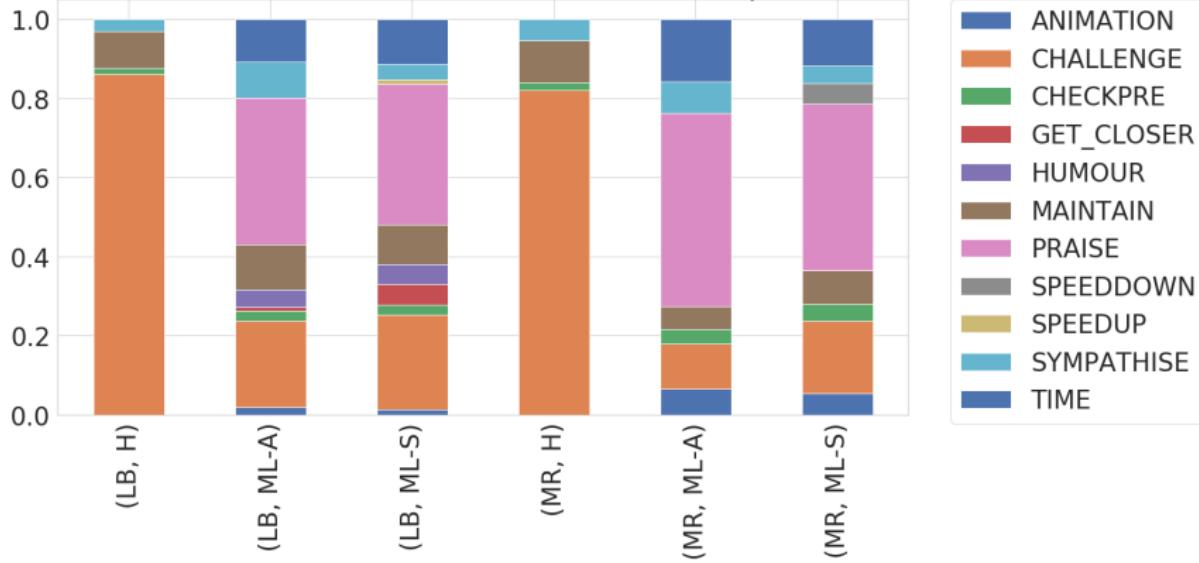


LEARNING TO IMITATE THE HUMAN EXPERT



LEARNING TO PERSONALISE

Phase 3 H, ML-A and ML-S Action Distribution for Participants LB and MR



Emotions?
oo

Classification
oooooooooooooooooooo

Emotion classifier
oooooooooooo

ML for social behaviours generation
oooooooooooo

IN THIS LECTURE

- o Basic classifiers: kNNs, SVM, MLP
- o Build your own emotion classifier from scratch!
- o ML to teach a robot social behaviours

Emotions?
oo

Classification
oooooooooooooooooooo

Emotion classifier
oooooooooooo

ML for social behaviours generation
oooooooooooo

That's all for today, folks!

Questions:

severin.lemaignan@brl.ac.uk

Slides:

github.com/severin-lemaignan/lecture-hri-ml-for-hri